

Home Assignment 6

Advanced Programming

qhz731

October 23, 2024

1 Introduction

In this assignment, I implemented a concurrent job scheduling system that supports multiple workers for the Stateful Planning Committee (SPC). The main functions of the system include adding and removing workers, job scheduling, job cancellation, time-out handling, and exception handling. The code uses the channel and GenServer patterns to implement message passing between SPC and workers, and each worker runs in an independent thread. For the convenience of testing, I have written a series of test cases that cover the main functions of the system. To run the tests, simply execute the `cabal test --test-show-details=always` command in the project root directory. Based on existing test cases, my solution is functionally correct, but there is still room for improvement in the robustness and error handling of the code in some boundary cases of worker management and concurrent job execution.

2 Task: Adding Workers

In this task, I implemented an SPC worker system based on message passing and state management. When SPC receives the `MsgWorkerAdd` message, it first checks whether the worker name already exists. If not, it creates a new worker process through `spawn` and manages the worker state with `WorkerM` monad. `WorkerM` tracks the currently executed job (`currentJob`) through `WorkerState`, achieving the encapsulation and management of the worker state. In terms of scheduling, SPC's `schedule` function identifies idle workers (by checking the records in `spcJobsRunning`) and assigns pending jobs to idle workers through the `workerIsIdle` function. During the job allocation process, the SPC status is updated (moving the job from pending to running), and a message to execute the job is sent to the worker. SPC has an independent `jobDoneChan` channel, which is used by workers to report job completion status to SPC. After receiving the job completion message, SPC will send `MsgJobDone` to `SPCServer` to call `jobDone` to update the job status, and the worker will idle again for subsequent job scheduling. The entire implementation adopts a hierarchical state management architecture. SPC is responsible for global scheduling and state maintenance, while the worker manages its own execution state through an independent state monad. The two realize interaction through the message passing mechanism.

3 Task: Job Cancellation

In this task, I implemented the job cancellation function. When SPC receives `MsgJobCancel`, it will first look for the target job and its corresponding worker in `spcJobsRunning`. If a worker is found to be executing the job, a cancellation request is sent to the worker through `MsgCancelJob` in `WorkerMsg`. When a worker executes a new job, it will create an independent execution thread through `forkIO` to run the job, and store the thread ID (`ThreadId`) together with the `JobId` in the `currentJob` field of `WorkerState`. This design makes it easy to terminate the execution thread of a specific job through `killThread` without affecting the operation of the worker's main thread. When processing the cancellation request, the worker will terminate the corresponding job thread, update its own status (set `currentJob` to `Nothing`), and report to SPC through `jobDoneChan` that the job ends in the `DoneCancelled` state. After receiving this message, SPC handles the subsequent process through the `jobDone` function.

4 Task: Timeouts

In this task, I adopted a centralized timeout management strategy to implement job timeout control. When the system assigns a job to a worker, it calculates the job deadline based on the `jobMaxSeconds` field of the job and records the information together with the job in `spcJobsRunning`. The SPC main thread triggers the `checkTimeouts` function through the `MsgTick` message received regularly (every second), and the `checkTimeouts` function obtains the current time and compares it with the deadlines of all running jobs. For jobs that exceed the time limit, SPC sends a `MsgCancelJob` message to the corresponding worker with the end reason `DoneTimeout`. This design reuses the same infrastructure as the job cancellation mechanism: after receiving the timeout cancellation request, the worker terminates the job execution thread (by calling `killThread` with the `ThreadId` stored in `WorkerState.currentJob`), updates its own status, and reports the job end to SPC through `jobDoneChan`. By centrally managing timeout detection at the SPC level, the system implements unified time management, avoids the complexity of distributed clock synchronization, and simplifies the implementation of workers, allowing them to focus on job execution without maintaining their own timers.

5 Task: Exceptions

In this task, when a worker receives a new job, it runs the job in an independent execution thread and wraps the job execution in an exception handling block. The Haskell exception handling mechanism is used to catch any exceptions (`SomeException`) that may occur when the worker executes the job. If the job is completed successfully, the worker reports the `Done` status to SPC through `jobDoneChan`; if an exception is caught, the `DoneCrashed` status is reported. This design ensures that the exception is confined to the job execution thread and does not spread to the worker's main thread or other parts of the SPC system, ensuring that the worker can recover after a job crashes and continue to process new jobs.

6 Task: Removing Workers

In this task, I implemented the worker's exit mechanism. The main challenge was in solving the state synchronization problem during the exit process. The implementation adopts a multilevel state processing strategy: First, the `workerStop` function sends the `MsgWorkerStop` message to the worker to trigger the exit process. After receiving the stop message, the worker checks the currentJob state in `WorkerState`. If there is a running job, it calls `killThread` to terminate the job thread and reports to SPC through `jobDoneChan` that the job ends with the `DoneCancelled` state.

On the worker side, the message processing function returns a bool value to indicate if the worker should stop. I used `workerLoop` instead of `forever`. When a stop message is received, it returns false, no longer calls `workerLoop` to terminate the loop, and then sends an exit notification to SPC through `workerGoneChan`. However, due to the asynchronous nature of the communication channel, the SPC may have a short state inconsistency time window after the worker is actually terminated, resulting in new jobs being incorrectly assigned to the terminated worker. To solve this problem, I implemented the `checkIfJobsIsRunningOnDeadWorkers` function to check and reset the job status assigned to the terminated worker before each message processing.

Although this design cannot completely avoid the existence of state inconsistency time windows, it ensures the final state consistency of the system through state checking and compensation mechanisms.

7 Q&A

7.1 What happens when a job is enqueued in an SPC instance with no workers, and a worker is added later? Which of your tests demonstrate this?

When a job is queued in an SPC instance with no workers, the job is added to the `spcJobsPending` list and remains in a waiting state. After adding the worker, the job will be assigned to the worker and executed. My test case "add-job" demonstrates this scenario: first, a new SPC instance is created, and a job is added. The job status is verified as `JobPending` at this time through `jobStatus`. Then a worker "worker1" is added. The system's schedule mechanism detects the idle worker and assigns the pending job to the worker for execution through `workerIsIdle`. At this time, check again that the job status has changed to `JobRunning`. After the job is executed, the status changes to `Done`, and the value of `ref` is successfully modified, verifying that the job was indeed executed. In order to further verify that the worker can continue to process jobs, the test case adds a second job and confirms its successful execution.

```
testCase "add-job" $ do
  spc <- startSPC
  ref <- newIORef False
  j <- jobAdd spc $
    Job (threadDelay 1000000 >> writeIORef ref True) 2
  r1 <- jobStatus spc j
```

```

r1 @?= JobPending
-- Add a worker
_ <- workerAdd spc "worker1"
r2 <- jobStatus spc j
r2 @?= JobRunning
r3 <- jobWait spc j
r3 @?= Done
v <- readIORef ref
v @?= True
-- Check if worker can execute other jobs
j2 <- jobAdd spc $ Job (writeIORef ref False) 1
r4 <- jobWait spc j2
r4 @?= Done
v2 <- readIORef ref
v2 @?= False,

```

7.2 Did you decide to implement timeouts centrally in SPC, or decentrally in the worker threads? What are the pros and cons of your approach? Is there any observable effect?

I chose to implement a centralized timeout management mechanism in SPC. When assigning jobs, SPC calculates the deadline of the job (current time plus jobMaxSeconds) and stores it in spcJobsRunning together with the job information. The checkTimeouts check is triggered by regular MsgTick messages, and a cancellation request is sent to the worker to cancel the timeout jobs.

Pros: 1. Centralized time management avoids the problem of distributed clock synchronization and ensures the accuracy of timeout determination. 2. Simplifies the implementation of workers. When implementing the worker I only need to focus on job execution without maintaining their own timers.

Cons: 1. Due to periodic checks (once per second), there may be a timeout determination delay of up to 1 second. 2. All timeout checks are performed in the SPC main thread, which may cause performance bottlenecks when the number of jobs is large. 3. If the SPC thread is blocked, it may affect the timeout detection of all jobs.

7.3 Which of your tests verify that if a worker executes a job that is canceled, times out, or crashes, then that worker can still be used to execute further jobs?

I wrote three test cases ("cancel-job", "timeout", "crash") to verify the availability of the worker after handling an abnormal job:

```

testCase "cancel-job" $ do
  spc <- startSPC
  ref <- newIORef False
  j <- jobAdd spc $

```

```

    Job (threadDelay 2000000 >> writeIORef ref True) 3
_ <- workerAdd spc "worker1"
jobCancel spc j
r <- jobWait spc j
r @?= DoneCancelled
v <- readIORef ref
v @?= False
-- Ensure the worker is still working
j2 <- jobAdd spc $ Job (writeIORef ref True) 1
r2 <- jobWait spc j2
r2 @?= Done
v2 <- readIORef ref
v2 @?= True,

```

These test cases cover three situations where the job is terminated abnormally, and verify the continued availability of the worker by submitting a new job after the abnormal job. The test results show that the worker is able to maintain normal working status after handling an abnormal job.

7.4 If a worker thread were somehow to crash (either due to a bug in the worker thread logic, or because a scoundrel `killThreads` it), how would that impact the functionality of the rest of SPC?

In my implementation, if the worker thread crashes unexpectedly, it will cause the SPC system state to be inconsistent. Since the worker fails to send the `MsgWorkerGone` message normally, the record of the worker in the `spcWorkers` list will not be removed, and the job being executed on the worker will continue to remain in the `spcJobsRunning` state until it is canceled due to timeout. Since SPC does not know that the worker has crashed, after the timed-out job is canceled, the SPC system will think that the worker is idle and continue to try to assign new jobs to the crashed worker.

This is indeed a defect in the implementation. Possible improvements include implementing a worker heartbeat mechanism and regularly checking whether the worker is alive.

8 Reference

I used the AI tool `Cursor` to generate part of the test cases. The generated tests are carefully checked to ensure the reliability.